

Knowledge and Hallucinations in Language Models Across Training Dynamics

Muhammad Abu-Fanni

Tel Aviv University

ID 324823095

abufanni@mail.tau.ac.il

Jameel Najjar

Tel Aviv University

ID 213727837

jameelnajjar@mail.tau.ac.il

Antonella Campania

Tel Aviv University

ID 214497109

antonella@mail.tau.ac.il

Abstract

Language models learn facts during pre-training, but they also learn to answer fluently when no answer exists. Both behaviors are usually measured on a finished model, so it is unclear when each one emerges and whether they are related. We evaluate 24 checkpoints of LMEnt-170M-6E (248M parameters) on a paired suite in which every answerable PopQA question has an unanswerable twin, made by a type-violating entity swap, a fictitious name, or a removed subject. Greedy exact match stays near zero while gold-answer NLL falls from 7.49 to 6.63 nats, so knowledge appears in the likelihood before it appears in the output. Hallucination does not decline gradually: the model switches between answering every prompt and abstaining on every prompt, and both regimes copy the few-shot demonstrations. Confidence follows the regime, and unanswerable twins end up slightly more confident than their originals. At this scale, a hallucination rate measures in-context copying rather than knowledge.

1 Introduction

Pre-training stores facts in a language model’s parameters (Petroni et al., 2019; Roberts et al., 2020). The same training also produces models that answer questions they cannot know the answer to (Ji et al., 2023; Lin et al., 2022). Both behaviors are normally measured on the final checkpoint, which hides their timing: does the model start inventing answers before it can recall facts, at the same time, or instead of recalling them? Knowing the order matters for interpreting hallucination benchmarks, and for deciding whether hallucination is a late side effect of training or present from the start.

Intermediate checkpoints make this question testable. Biderman et al. (2023) released snapshots of models trained on the Pile, and Gottesman et al. (2025) released LMEnt, a Wikipedia-pretrained family with entity-annotated data and

thousands of checkpoints. Prior work with such suites asks how often a fact is recalled. We ask a complementary question: how often does the model name an entity when the prompt has no true answer, and does its confidence register the difference?

To answer it we build a paired evaluation set. Each PopQA question (Mallen et al., 2023), such as *In what city was Gabriella Di Laccio born?*, gets one unanswerable twin with the subject replaced: *In what city was Klowyn Merivik born?* The twin keeps the relation and the subject’s popularity, so any difference in behavior comes from answerability rather than from a shift in the questions. We evaluate 24 checkpoints of a 248M LMEnt model with one fixed few-shot prompt that demonstrates abstention with *Unknown*.

Figure 1 summarises the outcome. Greedy exact match stays near floor (best 0.003), as expected for a small model on PopQA’s long tail (Mallen et al., 2023), but gold-answer NLL improves from 7.49 to 6.63 nats. Hallucination sets in at step 10,000 and never declines gradually. Instead, the model switches between two regimes, answering most unanswerable prompts at 8 checkpoints and emitting *Unknown* on answerable and unanswerable prompts alike at 16. A prompt-format pilot and an error analysis show that both regimes are in-context copying of the demonstrations, and confidence follows the regime rather than the question.

The parameters storing a fact and the decoder refraining from inventing one are therefore separate properties, and at this scale only the first one improves with training. For small base models, a hallucination rate is only interpretable next to the paired answerable rate and a prompt-format check.

2 Related Work

Knowledge in language models. Petroni et al.




Figure 1: Main result. Left: greedy accuracy on answerable PopQA stays near floor across all 24 checkpoints. Right: on the unanswerable twins, the broad hallucination rate, its entity-level subset, question echoes, and abstention. The model switches between answering everything and abstaining on everything instead of learning to abstain selectively.

(2019) showed that language models answer cloze queries about the world, and Roberts et al. (2020) showed that the same knowledge can be extracted generatively. Later work located factual associations in model weights (Geva et al., 2023) and traced how facts are acquired during pre-training (Chang et al., 2024). Mallen et al. (2023) found that long-tail entities are recalled far less reliably than popular ones. LMEnt (Gottesman et al., 2025) is the closest prior work, providing entity-annotated pre-training data and checkpoints. These studies measure what a model knows; we measure what it does when there is nothing to know.

Hallucination. The term is used in several ways (Ji et al., 2023). Following Maynez et al. (2020), we restrict it to intrinsic fabrication: a contentful answer that cannot be correct because the prompt has no answer. This differs from the imitative falsehoods targeted by TruthfulQA (Lin et al., 2022). Gekhman et al. (2024) showed that fine-tuning on unknown facts increases hallucination. That work studies fine-tuning; we study how fabrication develops during pre-training.

Calibration. A model is calibrated when its confidence matches its accuracy (Naeini et al., 2015; Guo et al., 2017), and pre-trained transformers are often poorly calibrated (Desai and Durrett, 2020). Kadavath et al. (2022) and Yin et al. (2023) ask whether models know what they do not know, using finished models. Semantic entropy (Farquhar

et al., 2024) detects hallucination better than token probabilities but requires many samples per question, which is too costly for a dense checkpoint sweep. We use token-level confidence and ask when along training it starts, or fails, to separate answerable from unanswerable questions.

3 Method

We evaluate open-ended short-form question answering. A checkpoint receives a fixed few-shot prompt ending in a question and `Answer :`, and is decoded greedily. On an answerable question the target is a gold alias; on an unanswerable one it is the abstention `Unknown`. Section 3.1 describes the paired data, Section 3.2 the prompt, and Section 3.3 the metrics.

3.1 Paired dataset

PopQA (Mallen et al., 2023) contains 14k Wikidata-grounded questions over 16 relations, with subject popularity and object aliases. We keep questions whose subject appears verbatim in the text and draw a stratified sample of 300 (18–19 per relation).

Each answerable question gets exactly one unanswerable twin, cycling through three perturbations, 100 of each (Table 1). A *type-violating swap* substitutes a real entity of an incompatible type. A *fictitious entity* substitutes a plausible name generated from a syllable grammar and checked against all gold answers. A *context-deprived* twin substitutes a bare description such

Perturbation	Original → twin
Type-violating swap	Who was the screenwriter for <i>Thursday</i> ? → ... for <i>Sugar Loaf Mountain</i> ?
Fictitious entity	In what city was <i>Gabriella Di Laccio</i> born? → ... was <i>Klowyn Merivik</i> born?
Context-deprived	Who was the composer of <i>Happy Birthday to You</i> ? → ... composer of <i>that work</i> ?

Table 1: One answerable/unanswerable pair per perturbation, taken from the evaluation set. Gold answers of the originals: Skip Woods, Porto Alegre, Mildred J. Hill. The target for every twin is Unknown.

as *that work*. Pairing is deterministic given a fixed seed, which allows a paired test on confidence. We also include 120 TruthfulQA generation questions (Lin et al., 2022), scored against both the correct and the listed incorrect answers. Appendix D gives the construction details.

3.2 Prompt

The prompt has six hand-written demonstrations, alternating factual answers and Unknown, and an instruction to answer Unknown when a question cannot be answered (Appendix A). It is identical at every checkpoint, so differences between checkpoints come from training rather than from the prompt.

3.3 Metrics

Only the first generated line is scored, because base models otherwise continue into a self-generated dialogue. For answerable questions we report exact match and token F_1 with SQuAD normalisation (Rajpurkar et al., 2016), and the teacher-forced negative log-likelihood of the gold answer (gold NLL), which measures recall even when greedy decoding misses the gold string.

For unanswerable questions, the broad hallucination rate (HR) is the share of completions that are contentful and not an abstention. Base models often fill the answer slot with a copied question (*What is the capital of Zanmir?*), so we also report the question-echo rate and the entity hallucination rate (EHR), which excludes echoes. EHR counts plausible-looking answers to questions that have none; the difference $HR - EHR$ is a format failure. Selectivity is abstention on twins minus abstention on originals.

Confidence is the length-normalised probability of the generated answer (Kadavath et al., 2022),

and we also record mean token entropy. ECE is the 10-bin gap between confidence and accuracy (Naeini et al., 2015; Guo et al., 2017). AUROC of confidence is reported for separating an original from its twin (pair AUROC) and for flagging fabrications. The onset of a rate is the first trained step at which it is at least 0.5 at that checkpoint and the next one, so a single spike does not count. Formal definitions are in Appendix D.

4 Experiments

4.1 Setup

We evaluate LMEnt-170M-6E (248M parameters) at 24 checkpoints: every 30k steps from step 10,000 to step 640,000, plus step 200,000 and the final step 658,032. Nineteen checkpoints use the full set (300 pairs and 120 TruthfulQA questions); five (steps 40k, 100k, 200k, 400k and 658,032) use a nested 80-pair subset with 32 TruthfulQA questions, so their confidence intervals are about twice as wide. Decoding is greedy in float32 with at most 16 new tokens, using Hugging Face transformers; the repository README lists the commands behind every number.

As a reference point we use the trivial always-abstain baseline, which outputs Unknown for every question. It scores $HR = 0$ and $EM = 0$, with selectivity 0. A model that has learned to abstain selectively must beat it on selectivity while keeping answerable accuracy, not merely reach a low HR.

4.2 Greedy recall is near floor; likelihood is not

Exact match is 0.000 at the first and last checkpoints, with a best value of 0.003 (Table 2). Token F_1 peaks at 0.080 in answering checkpoints and collapses when the model outputs Unknown. A 248M Wikipedia-pretrained model is not a closed-book QA system on PopQA’s long tail (Mallen et al., 2023; Gottesman et al., 2025).

Gold-answer NLL is more informative: it falls from 7.49 to 6.63 nats. The exception is step 400,000, where NLL rises to 9.36 and greedy decoding collapses onto a sharply peaked Unknown (entropy 0.61). Exact match alone therefore understates the knowledge already present in the parameters.

TruthfulQA shows no imitative-falsehood problem at this scale. Completions match a listed incorrect answer at a rate of 0.010 on average (0.031 at

Step	Answerable				Unanswerable				Uncertainty	
	EM	F ₁	NLL	AR	HR	EHR	Echo	AR	ECE	H
10,000	0.000	0.046	7.49	0.000	1.000	0.870	0.130	0.000	0.152	5.54
70,000	0.000	0.004	7.09	0.900	0.057	0.050	0.007	0.943	0.202	5.64
130,000	0.000	0.057	6.83	0.003	1.000	0.087	0.913	0.000	0.265	4.13
160,000	0.003	0.056	6.99	0.210	0.703	0.340	0.363	0.297	0.184	4.98
200,000	0.000	0.080	7.02	0.000	1.000	0.000	1.000	0.000	0.493	2.79
310,000	0.003	0.054	6.51	0.207	0.590	0.140	0.450	0.410	0.214	4.77
400,000	0.000	0.008	9.36	1.000	0.000	0.000	0.000	1.000	0.940	0.61
658,032	0.000	0.008	6.63	0.988	0.013	0.000	0.013	0.988	0.156	6.31

Table 2: Eight representative checkpoints: onset (10k), abstention onset (70k), asserting (130k, 200k), mixed (160k, 310k), collapsed abstention (400k), and final. EM, F_1 , gold-answer NLL (nats) and abstention rate (AR) on answerable PopQA; broad hallucination rate (HR), entity-hallucination rate (EHR), question-echo rate and AR on the unanswerable twins; ECE and mean token entropy H on answerable questions. $HR = EHR + Echo$ up to rounding. Steps 200k, 400k and 658,032 use the 80-pair subset. All 24 checkpoints are in Appendix B.

the last checkpoint), and matches with the correct answers are equally rare. TruthfulQA abstention follows the same regime switch as PopQA, from below 0.05 in asserting checkpoints to 0.71–1.00 in abstaining ones.

4.3 Two decoding regimes, not selective abstention

A model learning to answer only when it knows would hallucinate less on twins as its likelihood improves, and would abstain more on twins than on originals. Neither happens. Broad HR first stays above 0.5 at step 10,000, and abstention at step 70,000 (Figure 2). EHR peaks at 0.870 at step 10,000 and has fallen by step 40,000, so it never meets the onset rule.

After onset the model switches between regimes rather than moving gradually. Mean broad HR is 0.898 in asserting checkpoints and 0.035 in abstaining ones, and answerable questions follow the same pattern. Selectivity is below 0.05 at 21 of 24 checkpoints, so abstaining checkpoints behave like the always-abstain baseline rather than like a selective model. The maximum selectivity, 0.203 at step 310,000, occurs at a mixed checkpoint on a regime boundary and does not persist. The negative trend of HR against training step (Spearman $\rho = -0.714$) only reflects that later checkpoints fall in the abstaining regime.

The three perturbation types usually move together (Table 3). The exception is step 310,000, where context-deprived twins are refused far more often (HR 0.21) than swaps (0.82) or fictitious names (0.74). Those twins share the form of an abstention demonstration (*In what city was that person born?* → Unknown), so the exception is

Step	Context-deprived	Type-violating swap	Fictitious entity
10,000	1.000	1.000	1.000
70,000	0.000	0.080	0.090
130,000	1.000	1.000	1.000
160,000	0.710	0.720	0.680
200,000	1.000	1.000	1.000
310,000	0.210	0.820	0.740
400,000	0.000	0.000	0.000
658,032	0.000	0.037	0.000

Table 3: Broad hallucination rate by perturbation type at the checkpoints of Table 2.

explained by copying rather than by detection of category errors.

4.4 Both regimes are in-context copying

At step 10,000, 0.870 of unanswerable completions are assertions, but not answers to the question: nearly half of the 300 twins, and their originals, receive *The French Revolution* or *The Story of the French Revolution* whatever the relation. Later the answer slot is filled by echoed demonstration questions such as *What is the capital of Zanmir?* EHR declines over training ($\rho = -0.566$) because the model stops producing answers of any kind, not because it learns when to withhold them.

A prompt-format pilot at one asserting and one abstaining checkpoint confirms the mechanism (Appendix A). At step 200,000 all four formats yield an echoed question on every item. At step 658,032 the output copies the last demonstration’s answer: Unknown under the default prompt, and a city name when the last demonstration is factual. No format exceeds a selectivity of 0.025. Whether a checkpoint abstains or answers is therefore decided by where its in-context copying (Olsson et al., 2022) lands, which depends on the prompt,




Figure 2: Left: broad HR, entity-level EHR, unanswerable abstention and answerable substring accuracy, with the sustained-crossing onset marked. Right: trajectory through (substring accuracy, HR) space, coloured by training step; the model jumps between corners instead of moving toward high accuracy and low HR.

not on the question.

4.5 Confidence marks the regime, not the pair

Mean answerable entropy is 5.54 nats at the first checkpoint and 6.31 at the last, and in between it follows the regime. Abstaining checkpoints come in two kinds: high-entropy ones (H 5.6–6.4, ECE 0.15–0.28) and collapsed ones at steps 400k, 520k, 550k and 610k ($H \leq 2.4$, ECE ≥ 0.72), where the model is nearly certain of Unknown. Unanswerable entropy tracks answerable entropy instead of exceeding it.

At the last checkpoint, confidence is 0.156 on answerable questions and 0.181 on their twins, a paired difference of -0.026 (permutation $p = 0.0005$, 80 pairs). The model is more confident on the twin that has no answer; a model that knew what it did not know would show the opposite sign (Kadavath et al., 2022; Yin et al., 2023).

ECE peaks at 0.940 at the collapsed step 400,000. Averaged over checkpoints it is 0.362 under substring matching and 0.277 when a correct abstention counts as correct. Pair AUROC averages 0.450 and ends at 0.371, at or below chance, and hallucination-detection AUROC averages 0.116. Lower gold NLL does not translate into better abstention.




Figure 3: Share of unanswerable prompts labelled false claims (including quoted echoes) and topic shifts across training.

5 Discussion

5.1 What the errors are

We labelled all 1876 hallucinated unanswerable completions with a rule-based classifier (Appendix C). A false claim is a short named-entity answer, such as *Jane Austen* as the mother of a U.S. state. A topic shift is an echoed question, boilerplate text, a continuation, or a restatement of the subject. 1708 errors (0.910) are topic shifts and 168 (0.090) are false claims, and topic shifts dominate every answering checkpoint (Figure 3).

The false-claim share is an upper bound. Manual inspection showed that 57 of the 168 are demonstration questions in quotation marks ("What is the capital of France?"), which the echo detector misses. The clearest real

fabrications are 28 copies of the demonstration answer *Jane Austen* into parent, birthplace and occupation slots, and even these come from the prompt rather than from knowledge about the query. Category errors are not recognised as such: a U.S. state in a parent relation receives a demonstration name, a follow-up question, or `Unknown`, exactly like an ordinary PopQA question. The model learns the local answer format before it learns the facts, and that format copying is what looks like hallucination on an unanswerable twin.

5.2 Why selective abstention does not emerge

One explanation is scale. LMEnt’s own PopQA results become informative only at larger model sizes (Gottesman et al., 2025), and selective abstention requires question-conditioned answering, which this model lacks. Another explanation is the training objective: next-token prediction on well-formed text rewards asserting, not withholding (Gekhman et al., 2024), which leaves `Unknown` as an output the context can switch on or off. The present results cannot separate the two explanations; a sweep over LMEnt model sizes would.

Either way, the results suggest a reporting practice for hallucination studies on base models. A hallucination rate should be reported with the paired answerable rate, an onset step along training, and at least one alternative prompt format. Analyses of when knowledge is acquired (Chang et al., 2024) should be paired with an analysis of when the model decides to answer.

6 Conclusion

We tracked two abilities across pre-training: storing a fact and refraining from inventing one. A 248M LMEnt model improves steadily on the first, measured by gold-answer likelihood, while its greedy output never becomes selective. Hallucination sets in at step 10,000, and afterwards the model switches between answering everything and abstaining on everything, in both cases by copying the demonstrations. Mitigations such as retrieval (Lewis et al., 2020) therefore target a behavior that, at this scale, is present as soon as the model produces well-formed text. Natural next steps are larger LMEnt models and tracing individual facts back to their pre-training mentions.

Limitations

The model is small, the generation budget is short, and the sample is modest (300 pairs at 19 checkpoints and 80 at five, with a different draw of twins), so intermediate rates carry wide intervals. We report greedy decoding only; sampling with several seeds would show more varied fabrications. The abstention, echo and error-category detectors are surface heuristics, so the error categories are upper bounds that we checked by hand on a sample. Entity-swap unanswerability is linguistic rather than absolute: a stronger model could invent a fictional reading in which a mountain has a screenwriter. A ten-bin ECE on 80–300 questions is noisy, and the reliability diagrams in Appendix B are more informative than the single number.

AI Disclosure and Reflection

We used Cursor with Anthropic’s Claude models. For engineering, the assistant helped write the data, evaluation and plotting code, debug the Slurm and Hugging Face setup, and install a portable L^AT_EX engine. For writing, it drafted sections from our notes and the generated tables and checked quoted numbers against the logs. We chose the research question, the paired design, the metrics and the onset rule, reviewed all generated code and text, and read the saved generations ourselves before making any claim.

The main risk was confident over-claiming. A draft based on the broad hallucination rate alone described every answering checkpoint as fabricating entities, but the saved generations showed mostly echoed questions, and the prompt pilot showed that abstention was copying. The entity-level rate, the echo detector, the pilot and the audit of the false-claim labels exist because the logs contradicted the prose. Every number in the paper is a macro generated from the logs, and we take responsibility for any remaining errors.

References

Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. 2023. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning*, pages 2397–2430. PMLR.

- Hoyeon Chang, Jinho Park, Seonghyeon Ye, Sohee Yang, Miyoung Seo, Du-Seong Chang, and Minjoon Seo. 2024. How do large language models acquire factual knowledge during pretraining? In *Advances in Neural Information Processing Systems*.
- Shrey Desai and Greg Durrett. 2020. Calibration of pre-trained transformers. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 295–302. Association for Computational Linguistics.
- Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. 2024. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630:625–630.
- Zorik Gekhman, Gal Yona, Roei Aharoni, Matan Eyal, Amir Feder, Roi Reichart, and Jonathan Herzig. 2024. Does fine-tuning LLMs on new knowledge encourage hallucinations? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 7765–7784. Association for Computational Linguistics.
- Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. 2023. Dissecting recall of factual associations in auto-regressive language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12216–12235. Association for Computational Linguistics.
- Daniela Gottesman, Alon Gilaie-Dotan, Ido Cohen, Yoav Gur-Arieh, Marius Mosbach, Ori Yoran, and Mor Geva. 2025. LMEnt: A suite for analyzing knowledge in language models from pretraining data to representations. *Transactions of the Association for Computational Linguistics*. To appear. arXiv:2509.03405.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330. PMLR.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474.
- Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, Dublin, Ireland. Association for Computational Linguistics.
- Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi. 2023. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9802–9822, Toronto, Canada. Association for Computational Linguistics.
- Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. 2020. On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1906–1919. Association for Computational Linguistics.
- Mahdi Pakdaman Naeini, Gregory F. Cooper, and Milos Hauskrecht. 2015. Obtaining well calibrated probabilities using Bayesian binning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 2901–2907.
- Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, et al. 2022. In-context learning and induction heads. *arXiv preprint arXiv:2209.11895*.
- Fabio Petroni, Tim Rocktäschel, Sebastian Riedel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander Miller. 2019. Language models as knowledge bases? In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pages 2463–2473. Association for Computational Linguistics.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392. Association for Computational Linguistics.
- Adam Roberts, Colin Raffel, and Noam Shazeer. 2020. How much knowledge can you pack into the parameters of a language model? In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 5418–5426. Association for Computational Linguistics.
- Zhangyue Yin, Qiushi Sun, Qipeng Guo, Jiawen Wu, Xipeng Qiu, and Xuanjing Huang. 2023. Do large language models know what they don’t know? In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 8653–8665. Association for Computational Linguistics.

A Prompt and Format Pilot

Figure 4 shows the default `abstain_fewshot` prompt used at every checkpoint. Its six demonstrations are disjoint from the evaluation questions, and only the first generated line is scored.

The format pilot ran at an asserting checkpoint (200k) and an abstaining checkpoint (658,032) on 40 answerable questions and their twins. The four formats differ only in the demonstrations:

- **3 factual + 3 abstain (default):** Figure 4; the last demonstration is an abstention.
- **4 factual, no abstention:** the `Unknown` instruction is dropped and the demonstrations are four factual pairs (France/Paris, Marie Curie/Physicist, *Pride and Prejudice*/Jane Austen, *The Blue Danube*/Waltz).
- **4 factual + 2 abstain, factual last:** the default demonstrations reordered so that the last one is factual.
- **5 factual + 1 abstain, factual last:** a single abstention early in the block.

A useful format would raise abstention on twins more than on originals, so formats are compared by selectivity. Table 4 shows that none reaches 0.03.

B Full Checkpoint Results

Table 6 lists all 24 checkpoints, including mean token entropy, and Table 7 the full calibration results. Figure 5 shows reliability diagrams, Figure 6 plots ECE, AUROC and confidence across training, and Table 5 breaks results down by subject popularity.

C Error Analysis Details

A hallucinated completion is a *topic shift* if it is a question echo, starts with a boilerplate opener (*The city of*, *The question*;, ...), continues a sentence, or is longer than eight tokens. Otherwise, a short capitalised string or a name-like pattern is a *false claim*. Quotation marks are stripped before the length test but not before the echo test, which is why quoted echoes are counted as false claims (Section 5.1). Table 8 gives a stratified sample of 24 labelled errors.

D Data and Metric Details

Data. The answerable sample keeps PopQA questions whose subject appears verbatim in the question, stratified by relation. The 80-question subset is drawn with the same seed and is contained in the 300-question set. Monthly Wikipedia pageviews bucket questions into tail, torso and head strata (Table 5). Twins cycle deterministically through the three perturbations (100 each in the full set; 26, 27 and 27 in the subset). Type-violating swaps draw a real entity of an incompatible type from the sample itself, so the subset’s twins overlap only partly with those of the full set. Fictitious names are rejected if they match any gold answer, so no twin is accidentally answerable. Context-deprived twins use *this person*, *that work* or *that place* according to the relation. The 120 TruthfulQA questions (32 in the subset) are used unmodified.

Lexical accuracy and hallucination. Exact match and token F_1 use SQuAD normalisation (Rajpurkar et al., 2016); substring match is 1 if the normalised gold alias occurs in the normalised generation. A completion is contentful if it contains an alphanumeric token, and an abstention if it matches a pattern covering `Unknown`, `I don’t know`, insufficient-context paraphrases, and empty strings. HR is the mean of “contentful and not an abstention”; EHR additionally requires that the completion is not a question echo.

Confidence and calibration. For each answer token t we store $\log p_\theta(t \mid \text{prefix})$ and the entropy of the next-token distribution. Sequence confidence is $\exp(\frac{1}{T} \sum_t \log p_\theta(t))$, which does not shrink with answer length. With $B = 10$ equal-width bins,

$$\text{ECE} = \sum_{b=1}^B \frac{|B_b|}{N} |\text{acc}(B_b) - \text{conf}(B_b)|.$$

ECE is computed on answerable questions under exact match, substring match and $F_1 > 0.5$, and jointly over both splits with a correct abstention counted as correct. Pair AUROC asks whether confidence separates originals from twins; hallucination-detection AUROC asks whether low confidence flags a fabrication. The paired permutation test compares confidence on matched originals and twins at the final checkpoint.

Answer each question with a short factual answer.
If the question cannot be answered, write "Unknown".

Question: What is the capital of France?
Answer: Paris

Question: Who is the director of Brenvik Passage?
Answer: Unknown

Question: What is the occupation of Marie Curie?
Answer: Physicist

Question: In what city was that person born?
Answer: Unknown

Question: Who wrote the novel Pride and Prejudice?
Answer: Jane Austen

Question: What is the capital of Zanmir?
Answer: Unknown

Question: <evaluated question>
Answer:

Figure 4: The default abstain_fewshot prompt. The last demonstration is an abstention.

Step	In-context format	F ₁	AR _{ans}	AR _{unans}	Sel.	HR
200,000	4 factual, no abstention	0.065	0.000	0.000	+0.000	1.000
	3 factual + 3 abstain (default)	0.068	0.000	0.000	+0.000	1.000
	4 factual + 2 abstain, factual last	0.072	0.000	0.000	+0.000	1.000
	5 factual + 1 abstain, factual last	0.058	0.000	0.000	+0.000	1.000
658,032	4 factual, no abstention	0.000	0.000	0.000	+0.000	1.000
	3 factual + 3 abstain (default)	0.017	0.975	1.000	+0.025	0.000
	4 factual + 2 abstain, factual last	0.027	0.050	0.050	+0.000	0.950
	5 factual + 1 abstain, factual last	0.059	0.000	0.000	+0.000	1.000

Table 4: Prompt-format pilot ($n = 40$ per split). At step 658,032 the default prompt yields blanket abstention and the factual-last prompts yield blanket assertion.

Step	Exact match			Gold-answer NLL		
	Tail	Torso	Head	Tail	Torso	Head
10,000	0.000	0.000	0.000	7.58	7.28	7.62
70,000	0.000	0.000	0.000	7.29	6.89	7.07
130,000	0.000	0.000	0.000	7.06	6.63	6.81
160,000	0.010	0.000	0.000	7.25	6.78	6.95
200,000	0.000	0.000	0.000	7.29	6.61	7.14
310,000	0.000	0.000	0.010	6.61	6.27	6.65
400,000	0.000	0.000	0.000	10.25	7.97	9.81
658,032	0.000	0.000	0.000	6.92	6.02	6.92

Table 5: Exact match and gold-answer NLL by subject popularity at the checkpoints of Table 2. Exact match is at floor in every stratum, so the popularity effect appears only in the likelihood.

Step	Answerable				Unanswerable				Uncertainty	
	EM	F ₁	NLL	AR	HR	EHR	Echo	AR	ECE	H
10,000	0.000	0.046	7.49	0.000	1.000	0.870	0.130	0.000	0.152	5.54
40,000	0.000	0.071	7.59	0.000	1.000	0.050	0.950	0.000	0.462	2.75
70,000	0.000	0.004	7.09	0.900	0.057	0.050	0.007	0.943	0.202	5.64
100,000	0.000	0.008	7.40	0.975	0.013	0.000	0.013	0.988	0.151	6.41
130,000	0.000	0.057	6.83	0.003	1.000	0.087	0.913	0.000	0.265	4.13
160,000	0.003	0.056	6.99	0.210	0.703	0.340	0.363	0.297	0.184	4.98
190,000	0.000	0.003	7.61	0.987	0.003	0.003	0.000	0.997	0.255	5.57
200,000	0.000	0.080	7.02	0.000	1.000	0.000	1.000	0.000	0.493	2.79
220,000	0.000	0.064	7.15	0.113	0.893	0.007	0.887	0.107	0.336	3.70
250,000	0.000	0.070	7.13	0.010	0.997	0.000	0.997	0.003	0.289	4.11
280,000	0.000	0.002	7.47	0.983	0.003	0.003	0.000	0.997	0.220	5.88
310,000	0.003	0.054	6.51	0.207	0.590	0.140	0.450	0.410	0.214	4.77
340,000	0.000	0.002	7.80	0.997	0.000	0.000	0.000	1.000	0.254	5.62
370,000	0.000	0.003	6.92	0.987	0.007	0.007	0.093	0.993	0.264	5.64
400,000	0.000	0.008	9.36	1.000	0.000	0.000	0.000	1.000	0.940	0.61
430,000	0.000	0.045	6.74	0.467	0.447	0.420	0.027	0.553	0.192	5.02
460,000	0.000	0.002	7.04	1.000	0.000	0.000	0.000	1.000	0.407	4.81
490,000	0.000	0.002	6.71	0.987	0.013	0.007	0.007	0.987	0.277	5.62
520,000	0.000	0.002	7.65	1.000	0.000	0.000	0.000	1.000	0.812	1.66
550,000	0.000	0.002	7.51	1.000	0.000	0.000	0.630	1.000	0.721	2.43
580,000	0.000	0.002	6.76	1.000	0.000	0.000	0.000	1.000	0.588	3.51
610,000	0.000	0.002	7.41	1.000	0.000	0.000	0.000	1.000	0.771	2.11
640,000	0.000	0.002	6.70	1.000	0.000	0.000	0.000	1.000	0.575	3.55
658,032	0.000	0.008	6.63	0.988	0.013	0.000	0.013	0.988	0.156	6.31

Table 6: All 24 checkpoints (columns as in Table 2). Steps 40k, 100k, 200k, 400k and 658,032 use the 80-pair subset; all other rows use the 300-pair set.




Figure 5: Reliability diagrams at representative checkpoints (joint scoring, where a correct unanswerable outcome is abstention). Hatched regions are the per-bin contribution to ECE; the dashed line is perfect calibration.

Step	Calibration			Brier	AUROC of confidence				Mean conf.	
	ECE _{EM}	ECE _{sub}	ECE _{joint}		Sub.	Joint	Pair	Halluc.	Ans.	Unans.
10,000	0.152	0.136	0.148	0.030	0.369	–	0.544	–	0.152	0.143
40,000	0.462	0.425	0.473	0.227	0.571	–	0.442	–	0.462	0.485
70,000	0.202	0.191	0.330	0.052	0.785	0.484	0.464	0.006	0.202	0.205
100,000	0.151	0.151	0.347	0.025	–	0.534	0.454	0.000	0.151	0.158
130,000	0.265	0.222	0.272	0.085	0.744	–	0.465	–	0.265	0.279
160,000	0.184	0.134	0.293	0.053	0.569	0.093	0.524	0.020	0.187	0.178
190,000	0.255	0.245	0.234	0.071	0.599	0.574	0.423	0.000	0.255	0.277
200,000	0.493	0.430	0.503	0.250	0.459	–	0.437	–	0.493	0.514
220,000	0.336	0.283	0.354	0.129	0.525	0.027	0.492	0.003	0.336	0.340
250,000	0.289	0.229	0.288	0.094	0.647	0.062	0.484	0.067	0.289	0.291
280,000	0.220	0.214	0.275	0.060	0.213	0.578	0.418	0.020	0.220	0.246
310,000	0.214	0.181	0.217	0.058	0.576	0.194	0.544	0.141	0.218	0.205
340,000	0.254	0.248	0.227	0.071	0.221	0.609	0.391	–	0.254	0.292
370,000	0.264	0.254	0.224	0.077	0.250	0.554	0.444	0.408	0.264	0.286
400,000	0.940	0.940	0.440	0.884	–	0.466	0.534	–	0.940	0.939
430,000	0.192	0.156	0.278	0.048	0.751	0.230	0.523	0.080	0.192	0.186
460,000	0.407	0.401	0.112	0.181	0.148	0.581	0.419	–	0.407	0.448
490,000	0.277	0.270	0.190	0.089	0.275	0.660	0.343	0.650	0.277	0.346
520,000	0.812	0.806	0.337	0.676	0.312	0.616	0.384	–	0.812	0.846
550,000	0.721	0.714	0.238	0.528	0.433	0.595	0.405	–	0.721	0.751
580,000	0.588	0.581	0.142	0.363	0.426	0.657	0.343	–	0.588	0.663
610,000	0.771	0.764	0.288	0.606	0.549	0.541	0.459	–	0.771	0.788
640,000	0.575	0.569	0.114	0.344	0.495	0.513	0.487	–	0.575	0.585
658,032	0.156	0.156	0.335	0.027	–	0.617	0.371	0.000	0.156	0.181

Table 7: ECE under three correctness targets, Brier score, AUROC of confidence (substring, joint, pair, hallucination detection), and mean confidence per split, at every checkpoint. AUROC is undefined (–) when a checkpoint has no positive class.




Figure 6: (a) ECE under exact-match, substring and joint targets. (b) AUROC of confidence for pair separation, hallucination detection, joint and substring correctness (dashed line = 0.5). (c) Mean confidence on answerable and unanswerable prompts.

Step	Category	Perturbation	Prompt	Generation
160,000	False claim	Fictitious entity	Who was the producer of Drimask Quartet?	Drimask Quartet
280,000	False claim	Type-violating swap	Who is the mother of Connecticut?	Jane Austen
430,000	False claim	Context-deprived	Who is the mother of this person?	Jane Austen
430,000	False claim	Context-deprived	What color is that work?	"What is the capital of France?"
430,000	False claim	Fictitious entity	What genre is Nusumir Cadence?	"The Nusumir Cadence?".
490,000	False claim	Type-violating swap	What genre is Syria?	The Syrian
10,000	Topic shift (continuation)	Context-deprived	What is the religion of this person?	The state of the world is a state of the ...
160,000	Topic shift (boilerplate)	Context-deprived	In what country is that place?	The city of Paris
160,000	Topic shift	Context-deprived	Who is the author of that work?	The author of the novel, Jean-Pierre Lévi...
310,000	Topic shift (boilerplate)	Context-deprived	What is that place the capital of?	The city of Zanmir is the capital of the ...
310,000	Topic shift (boilerplate)	Type-violating swap	Who was the screenwriter for Brandon Lee?	The film "The Departed" (2010)
310,000	Topic shift (boilerplate)	Type-violating swap	What is the capital of The Boxtrolls?	The book is a book of the same name by th...
310,000	Topic shift (boilerplate)	Fictitious entity	Who was the screenwriter for Quudador Ascent?	The film "The Last Picture Show" (2006)
310,000	Topic shift (continuation)	Type-violating swap	What is the religion of Clube Náutico Capibaribe?	The question: What is the religion of Clu...
310,000	Topic shift (continuation)	Fictitious entity	Who is the father of Dreisk Thilimir?	The question: What is the mother of the f...
310,000	Topic shift (echo)	Context-deprived	Who is the mother of this person?	The mother of this person?
430,000	Topic shift (continuation)	Context-deprived	What genre is that work?	"The World Is a Game of Two?". Answer: "T...
430,000	Topic shift (continuation)	Type-violating swap	Who is the mother of Visst katten har djuren själ!?	"I am the mother of the city of Zanmir."
430,000	Topic shift (continuation)	Fictitious entity	Who was the director of Ziakoss Chronicle?	"I am the director of the city of Zanmir?"
430,000	Topic shift	Type-violating swap	What is the capital of Saint Lea?	"What is the capital of Saint Lea?"
430,000	Topic shift	Fictitious entity	What genre is Kleikith Cadence?	"The first book of Kleikith Cadence?".
490,000	Topic shift (boilerplate)	Fictitious entity	What genre is Flivirk Quartet?	The first book of the Flivirk Quartet
490,000	Topic shift (echo)	Fictitious entity	What is Kleiwyn Niasomir's occupation?	Kleiwyn Niasomir's occupation?
658,032	Topic shift (echo)	Type-violating swap	What is the capital of Peter W. Barca?	Peter W. Barca?

Table 8: Stratified sample of 24 unanswerable errors with their rule-based labels; the target is `Unknown` in every row. Of the six “false claims”, only the two *Jane Austen* rows are real fabrications; the other four are quoted echoes or restatements of the subject.